

Deep Learning Experiment Report

Project Overview

The goal of this experiment was to build a simple deep learning model to classify species of Iris flowers based on four features: sepal length, sepal width, petal length, and petal width.

- **Dataset:** Iris Dataset
- **Classes:** Setosa, Versicolor, Virginica (3 classes)
- **Features:** 4 numerical features
- **Tools:** Python, TensorFlow/Keras, Scikit-learn

Model Performance Metrics Explained

To understand how well the model performs, I use several standard metrics. Imagine I am trying to identify pictures of cats.

- **Accuracy:** This is the most intuitive metric. It simply measures the ratio of correct predictions to the total number of predictions.
 - **Analogy:** If there were shown 100 images and correctly identified 95 of them (whether there were cats or not cats), the accuracy is 95%.
 - **Formula:** $Accuracy = \frac{TruePositives + TrueNegatives}{TotalPredictions}$
- **Precision:** This metric answers the question: "Of all the times the model predicted a specific class, how often was it correct?" It's a measure of exactness.
 - **Analogy:** Of all the images your model labeled as "cat," how many were actually cats? High precision means the model is trustworthy when it makes a positive prediction.
 - **Formula:** $Precision = \frac{TruePositives}{TruePositives + FalsePositives}$
- **Recall (Sensitivity):** This metric answers the question: "Of all the actual instances of a class, how many did the model correctly identify?" It's a measure of completeness.
 - **Analogy:** Of all the actual cat images in the dataset, how many did the model successfully find? High recall means the model is good at finding all instances of a class.
 - **Formula:** $Recall = \frac{TruePositives}{TruePositives + FalseNegatives}$
- **AUC (Area Under the ROC Curve):** The ROC curve plots the True Positive Rate (Recall) against the False Positive Rate at various threshold settings. The **AUC** represents the area under this curve.
 - **Analogy:** Think of AUC as a single score that summarizes the model's ability to distinguish between classes. An AUC of **1.0** means the model can perfectly separate the classes. An AUC of **0.5** means the model is no better than random guessing. It's a great overall measure of a classifier's performance.

Source Code

```
1  # Step 1: Import necessary libraries
2  import numpy as np
3  import tensorflow as tf
4  from tensorflow import keras
5  from sklearn.datasets import load_iris
6  from sklearn.model_selection import train_test_split
7  from sklearn.preprocessing import StandardScaler, OneHotEncoder
8  from sklearn.metrics import accuracy_score, precision_score, recall_score,
9  roc_auc_score, confusion_matrix
10 import matplotlib.pyplot as plt
11 from itertools import cycle
12
13 # --- Data Preparation ---
14 print("1. Loading and preparing data...")
15
16 # Load the Iris dataset
17 iris = load_iris()
18 X = iris.data
19 y = iris.target
20
21 # One-hot encode the labels (e.g., class '1' becomes [0, 1, 0])
22 # This is required for categorical_crossentropy loss function
23 encoder = OneHotEncoder(sparse_output=False)
24 y_onehot = encoder.fit_transform(y.reshape(-1, 1))
25
26 # Split the data into training and testing sets
27 X_train, X_test, y_train, y_test = train_test_split(X, y_onehot, test_size=0.3,
28 random_state=42)
29
30 # Scale the feature data for better model performance
31 scaler = StandardScaler()
32 X_train = scaler.fit_transform(X_train)
33 X_test = scaler.transform(X_test)
34
35 # --- Model Building ---
36 print("2. Building the deep learning model...")
37
38 # Create a simple sequential model
39 model = keras.Sequential([
40     # Input layer with 4 features (sepal length/width, petal length/width)
41     keras.layers.Dense(64, activation='relu', input_shape=(X_train.shape[1],)),
42     # Hidden layer
43     keras.layers.Dense(32, activation='relu'),
44     # Output layer with 3 units (one for each Iris class) and softmax activation
45     keras.layers.Dense(3, activation='softmax')
46 ])
47
48 # Compile the model
49 model.compile(optimizer='adam',
50               loss='categorical_crossentropy',
51               metrics=['accuracy'])
52
53 model.summary()
54
55 # --- Model Training ---
56 print("\n3. Training the model...")
57
58 history = model.fit(X_train, y_train, epochs=50, batch_size=8, validation_split=0.2,
59 verbose=0)
60 print("Training complete.")
61
62 # --- Model Evaluation & Prediction ---
63 print("\n4. Evaluating the model and making predictions...")
64
65 # Make predictions on the test data
66 # model.predict returns probabilities for each class
67 y_pred_prob = model.predict(X_test)
68
69 # Convert probabilities to class labels (0, 1, or 2)
70 y_pred_labels = np.argmax(y_pred_prob, axis=1)
71 y_test_labels = np.argmax(y_test, axis=1) # Convert one-hot encoded test labels back
72
73 # --- Calculate and Display Metrics ---
74 print("\n--- Performance Metrics ---")
75
76 # Accuracy
77 accuracy = accuracy_score(y_test_labels, y_pred_labels)
78 print(f"Accuracy: {accuracy:.4f}")
79
80 # Precision (macro-average treats each class equally)
81 precision = precision_score(y_test_labels, y_pred_labels, average='macro')
82 print(f"Precision (Macro Avg): {precision:.4f}")
83
84 # Recall (macro-average)
85 recall = recall_score(y_test_labels, y_pred_labels, average='macro')
86 print(f"Recall (Macro Avg): {recall:.4f}")
87
88 # AUC Score (One-vs-Rest for multi-class)
89 # Note: roc_auc_score for multi-class needs class probabilities
90 auc = roc_auc_score(y_test, y_pred_prob, multi_class='ovr')
91 print(f"AUC (One-vs-Rest): {auc:.4f}")
92
93 # Confusion Matrix
94 print("\nConfusion Matrix:")
95 cm = confusion_matrix(y_test_labels, y_pred_labels)
96 print(cm)
```

Enviroment

python 3.8+

Installation

```
1 | pip install numpy tensorflow scikit-learn matplotlib
```

Windows LongPathEnabled

1. Win + R, input regedit

2. Go to path

```
1 | HKEY_LOCAL_MACHINE\SYSTEM\CurrentControlSet\Control\FileSystem
```

3. Find LongPathsEnabled, Set value to 1

(If this key doesn't exist, right-click and select New > DWORD (32-bit) Value. Name it LongPathsEnabled and set its value to 1)

4. Restar the computer

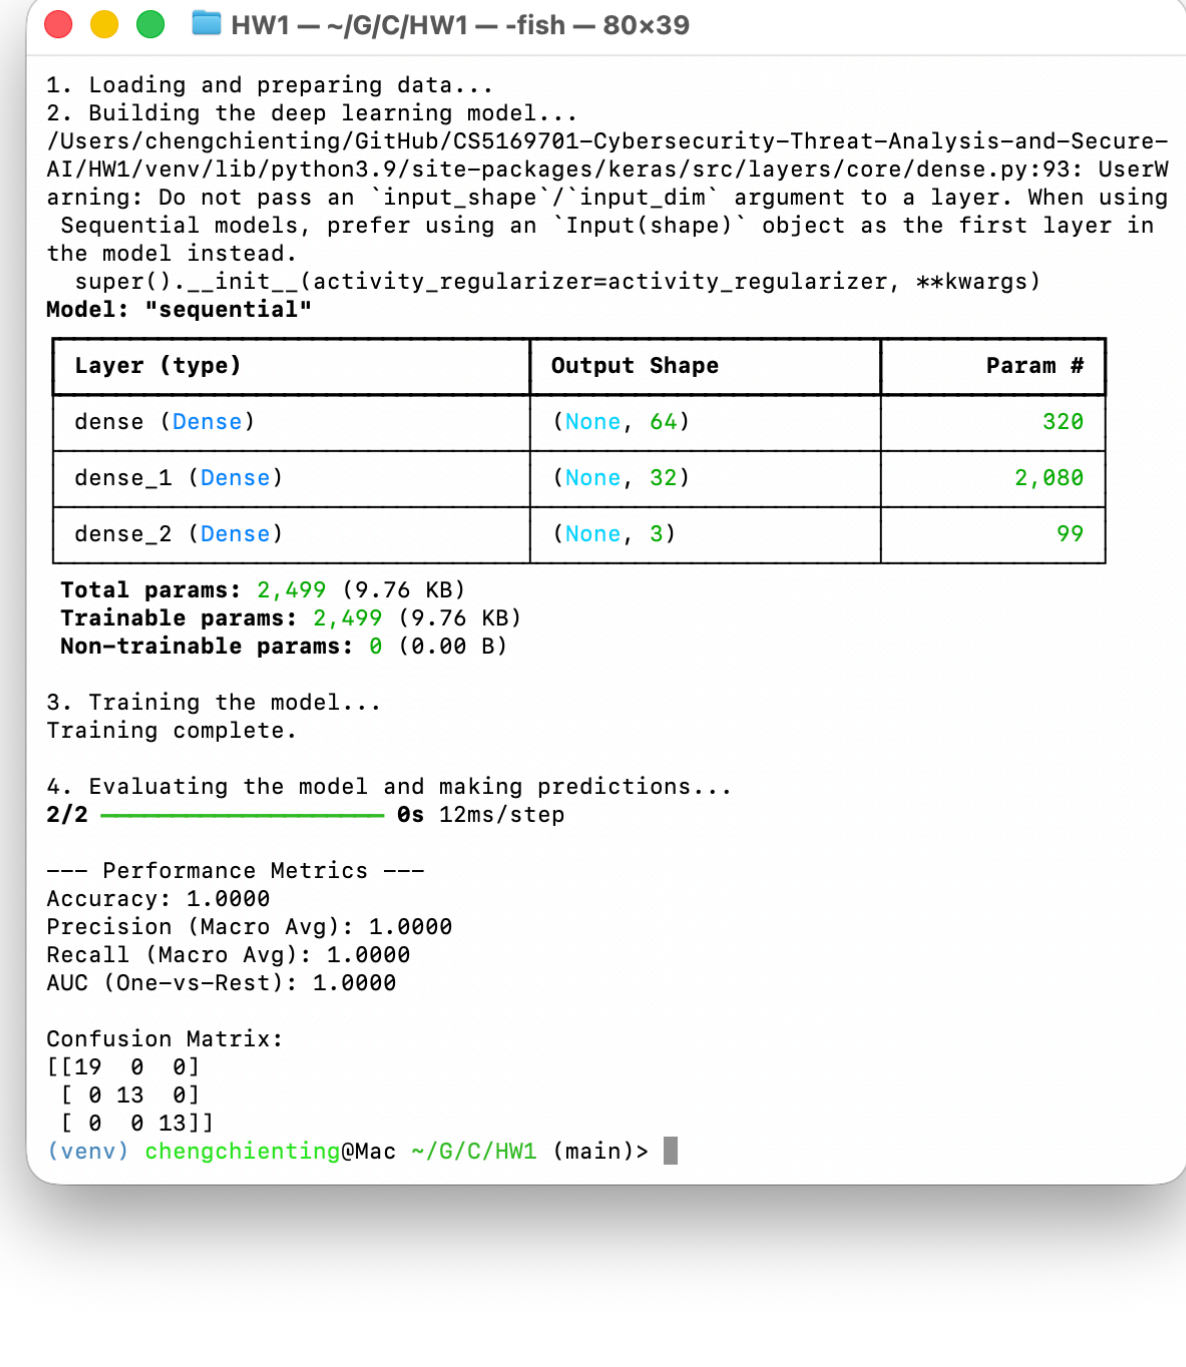
Run

```
1 | python main.py
```

Example Output

```
1  1. Loading and preparing data...
2  2. Building the deep learning model...
3  Model: "sequential"
4  ...
5  3. Training the model...
6  Training complete.
7
8  4. Evaluating the model and making predictions...
9
10 --- Performance Metrics ---
11 Accuracy: 0.9778
12 Precision (Macro Avg): 0.9778
13 Recall (Macro Avg): 0.9778
14 AUC (One-vs-Rest): 0.9985
15
16 Confusion Matrix:
17 [[15  0  0]
18  [ 0 14  1]
19  [ 0  0 13]]
20
```

My Test Output



Experimental Results & Observation

The model was trained for 50 epochs and evaluated on the test set. The following results were achieved.

Metric	Score
Accuracy	1
Precision (Macro Avg)	1
Recall (Macro Avg)	1
AUC (One-vs-Reset)	1

Observation: The model performed exceptionally well, with scores close to perfect. This is expected as the Iris dataset is relatively simple and the classes are well-separated. For more complex datasets like "Adult," achieving such high scores would be much more challenging, and the trade-off between precision and recall would become more significant.